

18-661 Introduction to Machine Learning

Nearest Neighbors

Fall 2020

ECE – Carnegie Mellon University

Midterm Information

Midterm will be on **Tuesday, 10/20 in-class**.

- Conducted as an online exam on Gradescope, with multiple-choice and short-answer questions
- Closed-book except for one double-sided letter-size handwritten page of notes that you can prepare as you wish.
- We will provide formulas for relevant probability distributions.
- You will not need a calculator. Only pen/pencil and scratch paper are allowed.

Will cover all topics up to and including Nearest Neighbors (10/15)

- (1) point estimation/MLE/MAP, (2) linear regression, (3) naive Bayes, (4) logistic regression, (5) SVMs, (6) Graphical Models, (7) Nearest Neighbors.
- Practice Midterm exam has been posted on Gradescope

1. Nearest Neighbor Classifier
2. Practical Aspects of NN

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs
 - Graphical Models

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs
 - Graphical Models

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs
 - Graphical Models
- Now we will discuss two *nonparametric* models:

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs
 - Graphical Models
- Now we will discuss two *nonparametric* models:
 - Nearest neighbors

Parametric vs. Nonparametric

- So far, we've discussed parametric machine learning models:
 - Linear regression
 - Naive Bayes
 - Logistic regression
 - Linear SVMs
 - Graphical Models
- Now we will discuss two *nonparametric* models:
 - Nearest neighbors
 - Decision trees

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
$$P(x|\theta, \mathcal{D}) = P(x|\theta).$$

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
 $P(x|\theta, \mathcal{D}) = P(x|\theta)$.
 - Often simpler and faster to learn, but can sometimes be a poor fit

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
 $P(x|\theta, \mathcal{D}) = P(x|\theta)$.
 - Often simpler and faster to learn, but can sometimes be a poor fit

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
$$P(x|\theta, \mathcal{D}) = P(x|\theta).$$
 - Often simpler and faster to learn, but can sometimes be a poor fit
- **Nonparametric models** instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
 $P(x|\theta, \mathcal{D}) = P(x|\theta)$.
 - Often simpler and faster to learn, but can sometimes be a poor fit
- **Nonparametric models** instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.
 - More complex and expensive, but can learn more flexible patterns

Parametric vs. Nonparametric

Key difference:

- **Parametric models** assume that the data can be characterized via some *fixed* set of parameters θ . Given this set of parameters, our future predictions are independent of the data \mathcal{D} , i.e.,
 $P(x|\theta, \mathcal{D}) = P(x|\theta)$.
 - Often simpler and faster to learn, but can sometimes be a poor fit
- **Nonparametric models** instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.
 - More complex and expensive, but can learn more flexible patterns
- Both parametric and non-parametric methods can be used for either regression or classification.

Nearest Neighbor Classifier


Recognizing flowers

Types of Iris: *setosa*, *versicolor*, and *virginica*



Measuring the properties of the flowers

Features: the widths and lengths of sepal and petal



Often, data is conveniently organized as a table

Ex: Iris data ([click here for all data](#))

- 4 features
- 3 classes

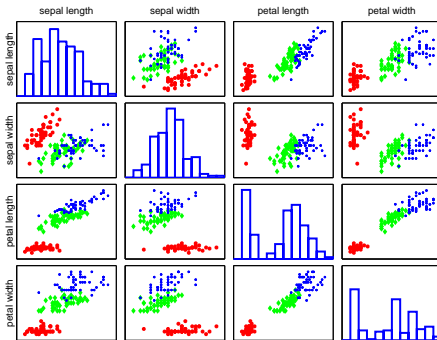
| Fisher's <i>Iris</i> Data | | | | |
|---------------------------|---------------|----------------|---------------|------------------|
| Sepal length ⇅ | Sepal width ⇅ | Petal length ⇅ | Petal width ⇅ | Species ⇅ |
| 5.1 | 3.5 | 1.4 | 0.2 | <i>I. setosa</i> |
| 4.9 | 3.0 | 1.4 | 0.2 | <i>I. setosa</i> |
| 4.7 | 3.2 | 1.3 | 0.2 | <i>I. setosa</i> |
| 4.6 | 3.1 | 1.5 | 0.2 | <i>I. setosa</i> |
| 5.0 | 3.6 | 1.4 | 0.2 | <i>I. setosa</i> |
| 5.4 | 3.9 | 1.7 | 0.4 | <i>I. setosa</i> |
| 4.6 | 3.4 | 1.4 | 0.3 | <i>I. setosa</i> |
| 5.0 | 3.4 | 1.5 | 0.2 | <i>I. setosa</i> |
| 4.4 | 2.9 | 1.4 | 0.2 | <i>I. setosa</i> |
| 4.9 | 3.1 | 1.5 | 0.1 | <i>I. setosa</i> |

Pairwise scatter plots of 131 flower specimens

Visualization of data helps to identify the right learning model

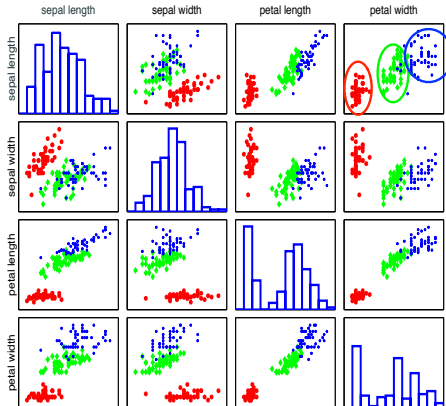
Which combination of features separates the three classes?

Figure 1: Each colored point is a flower specimen: **setosa**, **versicolor**, **virginica**

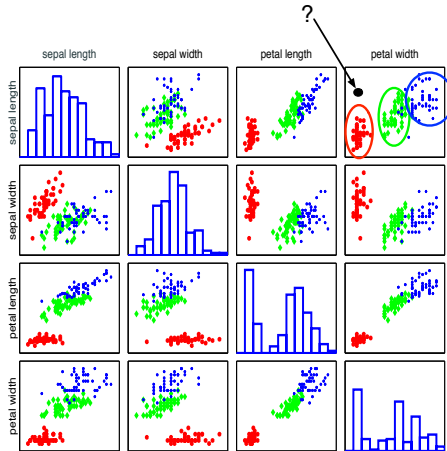


Different types seem well-clustered and separable

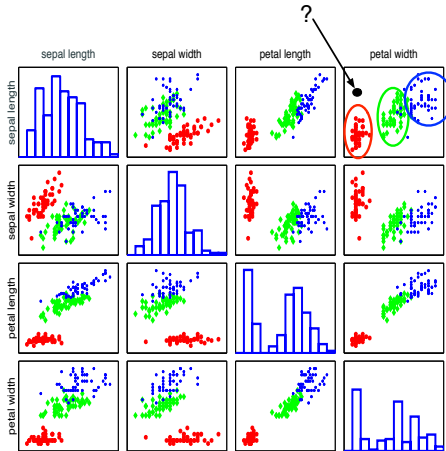
Using two features: petal width and sepal length



Labeling an unknown flower type



Labeling an unknown flower type



Closer to red cluster: so labeling it as **setosa**

Classify data into one of the multiple categories

- Input (feature vectors): $\mathbf{x} \in \mathbb{R}^D$
- Output (label): $y \in [C] = \{1, 2, \dots, C\}$
- Learning goal: $y = f(\mathbf{x})$

Classify data into one of the multiple categories

- Input (feature vectors): $\mathbf{x} \in \mathbb{R}^D$
- Output (label): $y \in [C] = \{1, 2, \dots, C\}$
- Learning goal: $y = f(\mathbf{x})$

Recall special case: binary classification

- Number of classes: $C = 2$
- Labels: $\{0, 1\}$ or $\{-1, +1\}$

Training data (set)

- N samples/instances: $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning $f(\cdot)$.

Training data (set)

- N samples/instances: $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning $f(\cdot)$.

Test (evaluation) data

- M samples/instances: $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well $f(\cdot)$ will do in predicting an unseen $\mathbf{x} \notin \mathcal{D}^{\text{TRAIN}}$.

Training data and test data should *not* overlap: $\mathcal{D}^{\text{TRAIN}} \cap \mathcal{D}^{\text{TEST}} = \emptyset$

Nearest neighbor classification (NNC)

Nearest neighbor of a (training or test) data point

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where $\text{nn}(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$, i.e., the index to one of the training instances

Nearest neighbor classification (NNC)

Nearest neighbor of a (training or test) data point

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where $\text{nn}(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$, i.e., the index to one of the training instances

$$\text{nn}(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2 = \operatorname{argmin}_{n \in [N]} \sum_{d=1}^D (x_d - x_{nd})^2$$

Nearest neighbor classification (NNC)

Nearest neighbor of a (training or test) data point

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where $\text{nn}(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$, i.e., the index to one of the training instances

$$\text{nn}(\mathbf{x}) = \underset{n \in [N]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2^2 = \underset{n \in [N]}{\operatorname{argmin}} \sum_{d=1}^D (x_d - x_{nd})^2$$

Classification rule

$$y = f(\mathbf{x}) = y_{\text{nn}(\mathbf{x})}$$

Nearest neighbor classification (NNC)

Nearest neighbor of a (training or test) data point

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where $\text{nn}(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$, i.e., the index to one of the training instances

$$\text{nn}(\mathbf{x}) = \underset{n \in [N]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2^2 = \underset{n \in [N]}{\operatorname{argmin}} \sum_{d=1}^D (x_d - x_{nd})^2$$

Classification rule

$$y = f(\mathbf{x}) = y_{\text{nn}(\mathbf{x})}$$


Example: if $\text{nn}(\mathbf{x}) = 2$, then

$$y_{\text{nn}(\mathbf{x})} = y_2,$$

which is the label of the 2nd data point.

Visual example

In this 2-dimensional example, the nearest point to \mathbf{x} is a **red training instance**, thus, \mathbf{x} will be labeled as **red**.



Example: classify Iris with two features

Training data

| ID (n) | petal width (x_1) | sepal length (x_2) | category (y) |
|------------|-----------------------|------------------------|------------------|
| 1 | 0.2 | 5.1 | setosa |
| 2 | 1.4 | 7.0 | versicolor |
| 3 | 2.5 | 6.7 | virginica |

Example: classify Iris with two features

Training data

| ID (n) | petal width (x_1) | sepal length (x_2) | category (y) |
|------------|-----------------------|------------------------|------------------|
| 1 | 0.2 | 5.1 | setosa |
| 2 | 1.4 | 7.0 | versicolor |
| 3 | 2.5 | 6.7 | virginica |

Flower with unknown category

petal width = 1.8 and sepal length = 6.4

Example: classify Iris with two features

Training data

| ID (n) | petal width (x_1) | sepal length (x_2) | category (y) |
|--------|-----------------------|------------------------|--------------|
| 1 | 0.2 | 5.1 | setosa |
| 2 | 1.4 | 7.0 | versicolor |
| 3 | 2.5 | 6.7 | virginica |

Flower with unknown category

petal width = 1.8 and sepal length = 6.4

Calculating distance from (x_1, x_2) to (x_{n1}, x_{n2}) : $(x_1 - x_{n1})^2 + (x_2 - x_{n2})^2$

| ID | distance |
|----|----------|
| 1 | 4.25 |
| 2 | 0.52 |
| 3 | 0.58 |

Thus, the predicted category is 2 (*versicolor*)

How to measure “nearness” with other distances?

Previously, we used Euclidean distance

$$\text{nn}(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$$


We can also use alternative distances

- E.g., the ℓ_1 distance (i.e., city block distance, or Manhattan distance):

$$\begin{aligned} \text{nn}(\mathbf{x}) &= \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_1 \\ &= \operatorname{argmin}_{n \in [N]} \sum_{d=1}^D |x_d - x_{nd}| \end{aligned}$$

- Or, the ℓ_∞ (supremum) distance:

$$\text{nn}(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \max_d |x_d - x_{nd}|$$



How to measure “nearness” with other distances?

Previously, we used Euclidean distance

$$\text{nn}(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$$

We can also use alternative distances

- E.g., the ℓ_1 distance (i.e., city block distance, or Manhattan distance):

$$\begin{aligned}\text{nn}(\mathbf{x}) &= \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_1 \\ &= \operatorname{argmin}_{n \in [N]} \sum_{d=1}^D |x_d - x_{nd}|\end{aligned}$$

- Or, the ℓ_∞ (supremum) distance:

$$\text{nn}(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \max_d |x_d - x_{nd}|$$






Figure 2: Green line is Euclidean distance. Red, Blue, and Yellow lines are L_1 distance

Decision boundary

For every point in the space, we can determine its label using the NNC rule. This gives rise to a *decision boundary* that partitions the space into different regions.




Recall: Multi-class classification



Previously, we learned a multi-class classifier by combining binary, **linear decision boundaries** to partition the feature space.

Recall: Multi-class classification



Previously, we learned a multi-class classifier by combining binary, **linear decision boundaries** to partition the feature space. Nearest neighbors can *naturally learn multiple decision boundaries* depending on which points are closest to which class's training data.

Nearest neighbors for regression

Recall the **nearest neighbor** of a (training or test) data point:

$$\mathbf{x}(1) = \mathbf{x}_{\text{nn}(\mathbf{x})}$$

where $\text{nn}(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$ indexes a training instance

$$\text{nn}(\mathbf{x}) = \underset{n \in [N]}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{x}_n\|_2^2 = \underset{n \in [N]}{\operatorname{argmin}} \sum_{d=1}^D (x_d - x_{nd})^2$$

Nearest neighbors for regression

Recall the **nearest neighbor** of a (training or test) data point:

$$\mathbf{x}(1) = \mathbf{x}_{nn(\mathbf{x})}$$


where $nn(\mathbf{x}) \in [N] = \{1, 2, \dots, N\}$ indexes a training instance

$$nn(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2 = \operatorname{argmin}_{n \in [N]} \sum_{d=1}^D (x_d - x_{nd})^2$$

Regression rule

$$y = f(\mathbf{x}) = y_{nn(\mathbf{x})}$$

Label \mathbf{x} with the label of its nearest neighbor!



Parametric vs. Nonparametric, revisited

Nonparametric models instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.

- Parametric models are often simpler and faster to learn, but can sometimes be a poor fit.
- Nonparametric models are **more complex and expensive**, but can **learn more flexible patterns**.

Parametric vs. Nonparametric, revisited

Nonparametric models instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.

- Parametric models are often simpler and faster to learn, but can sometimes be a poor fit.
- Nonparametric models are **more complex and expensive**, but can **learn more flexible patterns**.

How does this manifest for nearest neighbors?

- Nearest neighbors often learns a *highly nonlinear* decision boundary.

Parametric vs. Nonparametric, revisited

Nonparametric models instead assume that the model features depend on the data \mathcal{D} . The number of features tends to grow with the size of the dataset.

- Parametric models are often simpler and faster to learn, but can sometimes be a poor fit.
- Nonparametric models are **more complex and expensive**, but can **learn more flexible patterns**.

How does this manifest for nearest neighbors?

- Nearest neighbors often learns a *highly nonlinear* decision boundary.
- But, we need to compare the test data point to *every sample in the training dataset*.

K-nearest neighbor (KNN) classification

Increase the number of nearest neighbors to use?

- 1-nearest neighbor: $nn_1(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$

K-nearest neighbor (KNN) classification

Increase the number of nearest neighbors to use?

- 1-nearest neighbor: $nn_1(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 2nd-nearest neighbor: $nn_2(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$

K-nearest neighbor (KNN) classification

Increase the number of nearest neighbors to use?

- 1-nearest neighbor: $nn_1(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 2nd-nearest neighbor: $nn_2(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 3rd-nearest neighbor: $nn_3(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x}) - nn_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$

K-nearest neighbor (KNN) classification

Increase the number of nearest neighbors to use?

- 1-nearest neighbor: $nn_1(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 2nd-nearest neighbor: $nn_2(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 3rd-nearest neighbor: $nn_3(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x}) - nn_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$

K-nearest neighbor (KNN) classification

Increase the number of nearest neighbors to use?

- 1-nearest neighbor: $nn_1(\mathbf{x}) = \operatorname{argmin}_{n \in [N]} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 2nd-nearest neighbor: $nn_2(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$
- 3rd-nearest neighbor: $nn_3(\mathbf{x}) = \operatorname{argmin}_{n \in [N] - nn_1(\mathbf{x}) - nn_2(\mathbf{x})} \|\mathbf{x} - \mathbf{x}_n\|_2^2$

The set of K-nearest neighbors

$$\text{knn}(\mathbf{x}) = \{nn_1(\mathbf{x}), nn_2(\mathbf{x}), \dots, nn_K(\mathbf{x})\}$$

Let $\mathbf{x}(k) = \mathbf{x}_{nn_k(\mathbf{x})}$, then

$$\|\mathbf{x} - \mathbf{x}(1)\|_2^2 \leq \|\mathbf{x} - \mathbf{x}(2)\|_2^2 \leq \dots \leq \|\mathbf{x} - \mathbf{x}(K)\|_2^2$$

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then
 - vote for c is 1

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then
 - vote for c is 1
 - vote for $c' \neq c$ is 0

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then
 - vote for c is 1
 - vote for $c' \neq c$ is 0

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then
 - vote for c is 1
 - vote for $c' \neq c$ is 0

We use the *indicator function* $\mathbb{I}(y_n == c)$ to represent the votes.

- Aggregate everyone's vote

$$v_c = \sum_{n \in \text{knn}(\mathbf{x})} \mathbb{I}(y_n == c), \quad \forall \quad c \in [C]$$

How to classify with K neighbors?

Classification rule

- Every neighbor votes: suppose y_n (the true label) for \mathbf{x}_n is c , then
 - vote for c is 1
 - vote for $c' \neq c$ is 0

We use the *indicator function* $\mathbb{I}(y_n == c)$ to represent the votes.




- Aggregate everyone's vote

$$v_c = \sum_{n \in \text{knn}(\mathbf{x})} \mathbb{I}(y_n == c), \quad \forall \quad c \in [C]$$


- Label with the majority, breaking ties arbitrarily

$$y = f(\mathbf{x}) = \arg \max_{c \in [C]} v_c$$



Example






Example






Example



Example



How to choose an optimal K ?



When K increases, the decision boundary becomes smooth.

How to do regression with K neighbors?

- We need a way to aggregate labels from each of the neighbors.

How to do regression with K neighbors?

- We need a way to aggregate labels from each of the neighbors.
- **Average** the labels associated with the K points.

$$\hat{y} = \frac{1}{K} \sum_{n \in knn(\mathbf{x})} y_n$$

How to do regression with K neighbors?


- We need a way to aggregate labels from each of the neighbors.
- **Average** the labels associated with the K points.

$$\hat{y} = \frac{1}{K} \sum_{n \in knn(\mathbf{x})} y_n$$

How to do regression with K neighbors?

- We need a way to aggregate labels from each of the neighbors.
- **Average** the labels associated with the K points.

$$\hat{y} = \frac{1}{K} \sum_{n \in knn(\mathbf{x})} y_n$$



Why use nearest neighbors?

Advantages of NNC

- Computationally, simple and easy to implement – just compute distances, no optimization required
- Can learn complex decision boundaries

Why use nearest neighbors?

Advantages of NNC

- Computationally, simple and easy to implement – just compute distances, no optimization required
- Can learn complex decision boundaries

Disadvantages of NNC

- Computationally intensive for large-scale problems: $O(ND)$ for **labeling** a data point.
- We need to “carry” the training data around. Without it, we cannot do classification. This type of method is called *nonparametric*.
- Choosing the right distance measure and K can be difficult.

When should we not use nearest neighbors?

- Relies on the existence of training data points “close” to test points.
Choose your distance wisely!

When should we not use nearest neighbors?

- Relies on the existence of training data points “close” to test points. Choose your distance wisely!
- Can break down if the classes are unbalanced.

When should we not use nearest neighbors?

- Relies on the existence of training data points “close” to test points. Choose your distance wisely!
- Can break down if the classes are unbalanced.
- Can break down if we have irrelevant features. Solution—feature selection!

When should we not use nearest neighbors?

- Relies on the existence of training data points “close” to test points. Choose your distance wisely!
- Can break down if the classes are unbalanced.
- Can break down if we have irrelevant features. Solution—feature selection!
 - Essentially, we want to measure the “information” contributed by each feature.

When should we not use nearest neighbors?

- Relies on the existence of training data points “close” to test points. Choose your distance wisely!
- Can break down if the classes are unbalanced.
- Can break down if we have irrelevant features. Solution—feature selection!
 - Essentially, we want to measure the “information” contributed by each feature.
 - More on this in the next lecture on decision trees.

Practical Aspects of NN

Two crucial choices for NN

- Choosing K , i.e., the number of nearest neighbors (default is 1)

Two crucial choices for NN

- Choosing K , i.e., the number of nearest neighbors (default is 1)
- Choosing the right distance measure (default is Euclidean distance), for example, from the following generalized distance measure

$$\|\mathbf{x} - \mathbf{x}_n\|_p = \left(\sum_d |x_d - x_{nd}|^p \right)^{1/p}$$

for $p \geq 1$.

Two crucial choices for NN

- Choosing K , i.e., the number of nearest neighbors (default is 1)
- Choosing the right distance measure (default is Euclidean distance), for example, from the following generalized distance measure

$$\|\mathbf{x} - \mathbf{x}_n\|_p = \left(\sum_d |x_d - x_{nd}|^p \right)^{1/p}$$

for $p \geq 1$.

Two crucial choices for NN

- Choosing K , i.e., the number of nearest neighbors (default is 1)
- Choosing the right distance measure (default is Euclidean distance), for example, from the following generalized distance measure

$$\|\mathbf{x} - \mathbf{x}_n\|_p = \left(\sum_d |x_d - x_{nd}|^p \right)^{1/p}$$

for $p \geq 1$.

These are not specified by the algorithm itself — resolving them requires empirical studies and are task/dataset-specific.

Hyperparameter tuning on a validation dataset

Training data

- N samples/instances: $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning $f(\cdot)$

Hyperparameter tuning on a validation dataset

Training data

- N samples/instances: $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning $f(\cdot)$

Test data

- M samples/instances: $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well $f(\cdot)$ will do in predicting an unseen $\mathbf{x} \notin \mathcal{D}^{\text{TRAIN}}$

Hyperparameter tuning on a validation dataset

Training data

- N samples/instances: $\mathcal{D}^{\text{TRAIN}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- They are used for learning $f(\cdot)$

Test data

- M samples/instances: $\mathcal{D}^{\text{TEST}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_M, y_M)\}$
- They are used for assessing how well $f(\cdot)$ will do in predicting an unseen $\mathbf{x} \notin \mathcal{D}^{\text{TRAIN}}$

Validation data

- L samples/instances: $\mathcal{D}^{\text{VAL}} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_L, y_L)\}$
- They are used to optimize hyperparameter(s).

Training data, validation and test data should *not* overlap!

- For each possible value of the hyperparameter (say $K = 1, 3, \dots, 100$)

- For each possible value of the hyperparameter (say $K = 1, 3, \dots, 100$)
 - Train a model using $\mathcal{D}^{\text{TRAIN}}$ (we don't need this step for NNC)

- For each possible value of the hyperparameter (say $K = 1, 3, \dots, 100$)
 - Train a model using $\mathcal{D}^{\text{TRAIN}}$ (we don't need this step for NNC)
 - Evaluate the performance of the model on \mathcal{D}^{VAL}

- For each possible value of the hyperparameter (say $K = 1, 3, \dots, 100$)
 - Train a model using $\mathcal{D}^{\text{TRAIN}}$ (we don't need this step for NNC)
 - Evaluate the performance of the model on \mathcal{D}^{VAL}
- Choose the model with the best performance on \mathcal{D}^{VAL}


- For each possible value of the hyperparameter (say $K = 1, 3, \dots, 100$)
 - Train a model using $\mathcal{D}^{\text{TRAIN}}$ (we don't need this step for NNC)
 - Evaluate the performance of the model on \mathcal{D}^{VAL}
- Choose the model with the best performance on \mathcal{D}^{VAL}
- Evaluate the model on $\mathcal{D}^{\text{TEST}}$

Cross-validation

What if we do not have validation data?

- We split the training data into S equal parts.
- We use each part *in turn* as a validation dataset and use the others as a training dataset.
- We choose the hyperparameter such that the model performs the best (based on average, variance, etc.)
- We re-train the model on the full training dataset with the best hyperparameter.

Figure 3: $S = 5$: 5-fold cross validation




Cross-validation

What if we do not have validation data?

- We split the training data into S equal parts.
- We use each part *in turn* as a validation dataset and use the others as a training dataset.
- We choose the hyperparameter such that the model performs the best (based on average, variance, etc.)
- We re-train the model on the full training dataset with the best hyperparameter.

Special case: when $S = N$, this will be leave-one-out.

Figure 3: $S = 5$: 5-fold cross validation



But how do we choose the distances?

Distances depend on units of the features!

Normalize data to have zero mean and unit standard deviation in each dimension

- Compute the means and standard deviations in each feature

$$\bar{x}_d = \frac{1}{N} \sum_n x_{nd}, \quad s_d^2 = \frac{1}{N-1} \sum_n (x_{nd} - \bar{x}_d)^2$$

- Scale the feature accordingly

$$x_{nd} \leftarrow \frac{x_{nd} - \bar{x}_d}{s_d}$$

Normalize data to have zero mean and unit standard deviation in each dimension

- Compute the means and standard deviations in each feature

$$\bar{x}_d = \frac{1}{N} \sum_n x_{nd}, \quad s_d^2 = \frac{1}{N-1} \sum_n (x_{nd} - \bar{x}_d)^2$$

- Scale the feature accordingly

$$x_{nd} \leftarrow \frac{x_{nd} - \bar{x}_d}{s_d}$$

Many other ways of normalizing data — you would need/want to try different ones and pick among them using (cross) validation

Summary so far

- Described a simple *nonparametric* learning algorithm
- Discussed a few practical aspects, such as tuning hyperparameters, with cross-validation

Good luck on the midterm! Please come to office hours and/or the recitation with questions.